



DMA over PCIe with FPGA.

Sven Stubbe, Jan Marjanović (both DESY)

2019-12-03

MTCA Pre-Workshop 2019

microTca
TECHNOLOGY LAB

HELMHOLTZ
RESEARCH FOR GRAND CHALLENGES



The objective of this tutorial session is to build a simple data acquisition system, in which the data is moved from an external device through a FPGA over PCIe into a CPU memory. There the data will be made available to user applications. We will use the DAMC-TCK7, its Board Support Package and a MicroTCA crate as a starting point. Then we will add a couple of Xilinx IP blocks and use Xilinx Linux driver to interface with FPGA modules. Our final goal is to have the data available for getting operated using a small Python application.

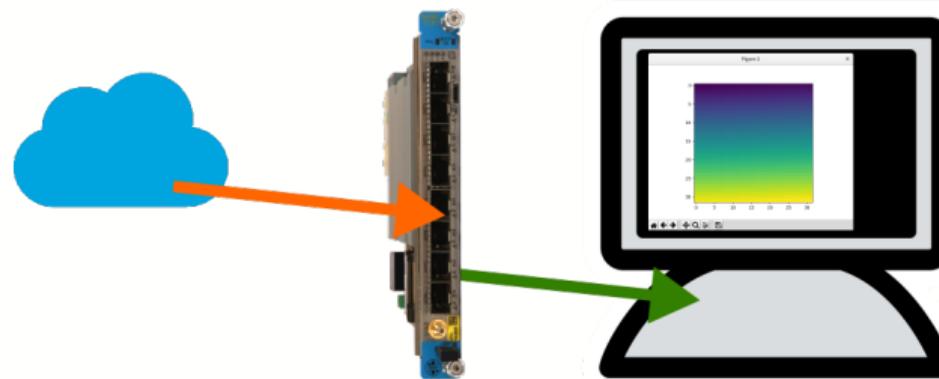
Speaker: Sven Stubbe, MicroTCA Technology Lab

Date: Tuesday, 3 December 2019 13:30-15:00

Venue: CFEL, SR I-III

Motivation for this talk:

Student with a task to transfer data from a detector to his software/Matlab/...

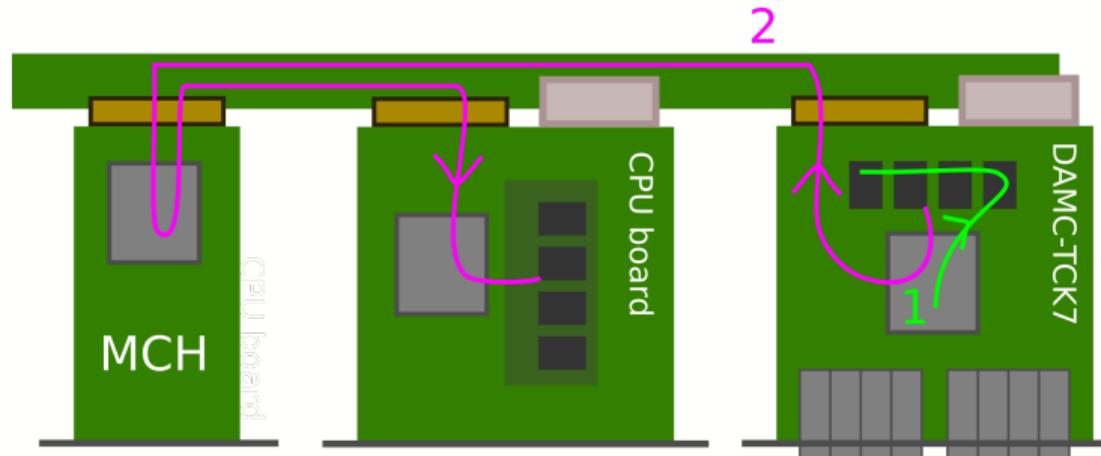


From Wikipedia:

Direct memory access (DMA) is a feature of computer systems that allows certain hardware subsystems [this is us] to access main system memory (random-access memory), independent of the central processing unit (CPU).

What is Direct Memory Access (DMA)?

Real example: move the data from a PCIe board into CPU memory

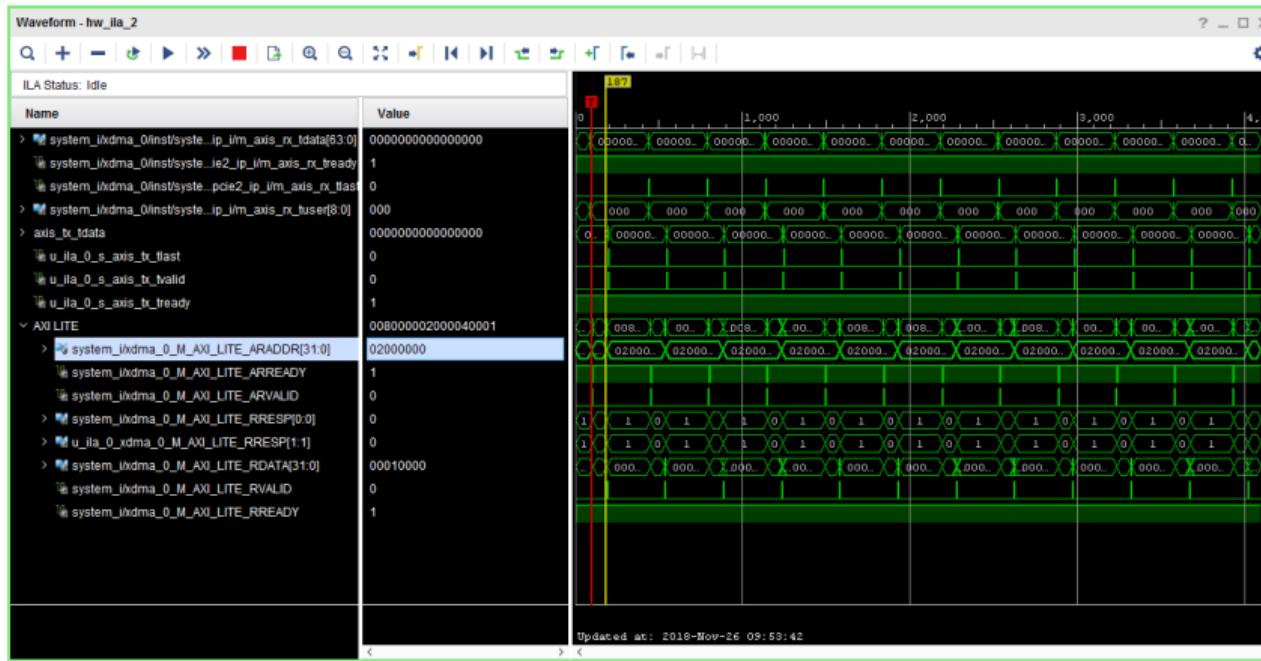


Direct memory access (DMA) is a feature of computer systems that allows certain hardware subsystems to access main system memory (random-access memory), independent of the central processing unit (CPU).

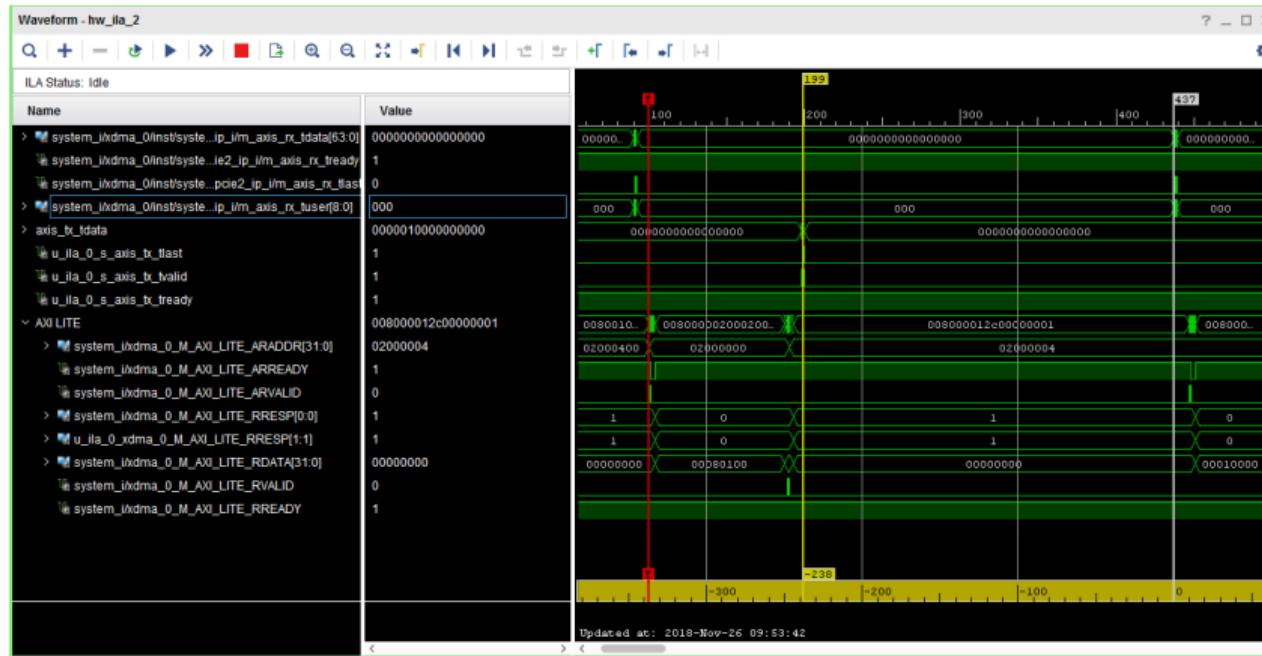
When we need to transfer large amounts of data quickly.

Alternative to DMA transfers is some kind of `memcpy()`:
processor reads the data into its registers, and writes the data out of registers. The CPU
manages the data transfers;
several CPU cycles for each data word are needed

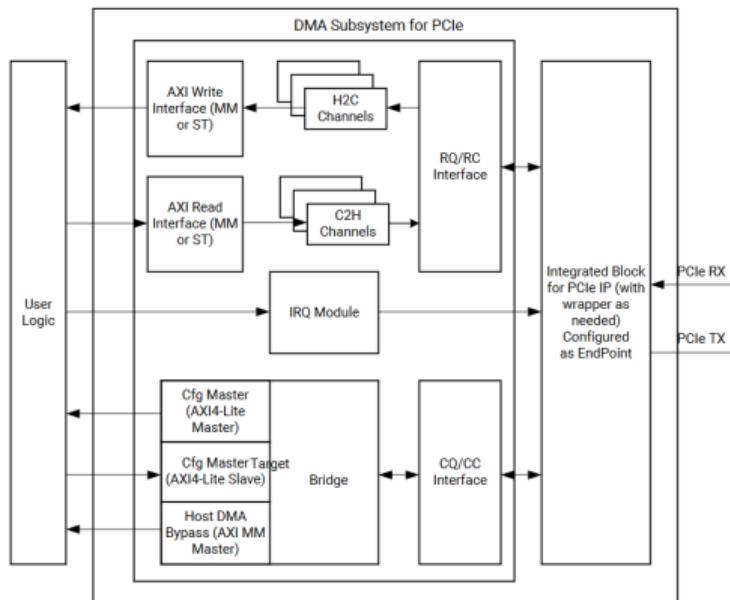
`memcpy()` over the PCIe in Linux kernel space results in ~ 2 MB/s transfer rate



A large delay between new MRd request (@437) and completion (@199) can be noted; this is not under our control



The CPU let the endpoint manage the the data transfers on its own:



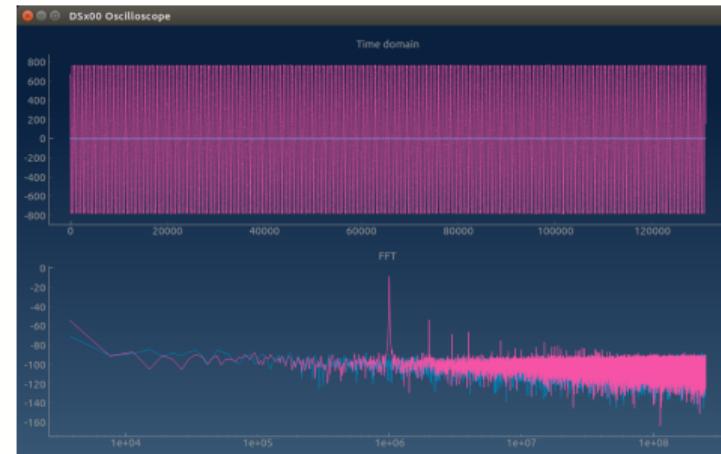
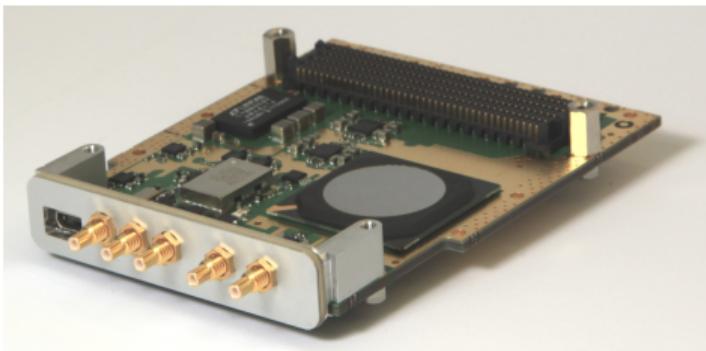
DMA/Bridge Subsystem for PCI Express® Overview (from pg195)

Examples

Example 1: DAQ for DFMC-DS800

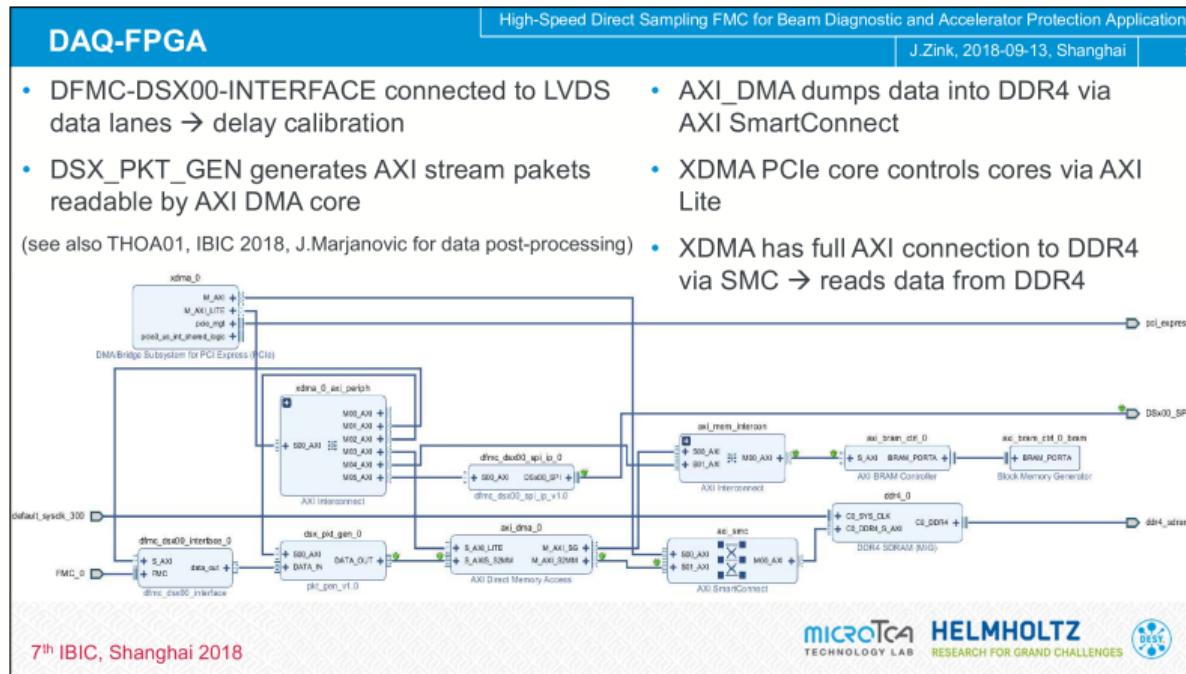
DMA over PCIe with FPGA
Stubbe, Marjanović, 2019-12-03 Page 11/70

- ▶ DFMC-DSx00 is an FMC mezzanine card with two-channel, 800 MSPS, 12-bit ADC
- ▶ Data rate is very high: 2400 MB/s
- ▶ DAQ application on Xilinx KCU105 eval kit



Example 1: DAQ for DFMC-DS800

DMA over PCIe with FPGA
Stubbe, Marjanović, 2019-12-03 Page 12/70



from http://ibic2018.vrws.de/talks/thoa02_talk.pdf

DAQ-Software	High-Speed Direct Sampling FMC for Beam Diagnostic and Accelerator Protection Applications J.Zink, 2018-09-13, Shanghai 6
<ul style="list-style-type: none">straight forward solution for DAQXDMA transfers sample data from DDR4 (KCU105) into host PC's RAM via PCIe root complex (DMA controller)Xilinx DMA driver (kernel space) provides pointer for DMA transfers into RAMPython library DSX00Lib can access transferred data in kernel space	<ul style="list-style-type: none">data can be dumped into file on SSD/HDDover 1 Mio. samples per channel can be stored on KCU105 (2GB DDR4)8k samples can be stored on DAMC-FMC25 (256 MB DDR2)

KCU105
FPGA
XDMA PCIe core

HOST PC
RAM
CPU
PCIe Root Complex

Xilinx DMA Driver (kernel space)
Python DSX00Lib (user space)
SSD (file system)
file

The diagram illustrates the data flow. Data from the KCU105 FPGA's XDMA PCIe core is transferred via the PCIe Root Complex to the HOST PC's RAM. This data is then processed by the CPU and the Xilinx DMA Driver (running in kernel space). The driver interacts with the Python DSX00Lib (running in user space), which finally writes the data to an SSD (file system) as a specific file.

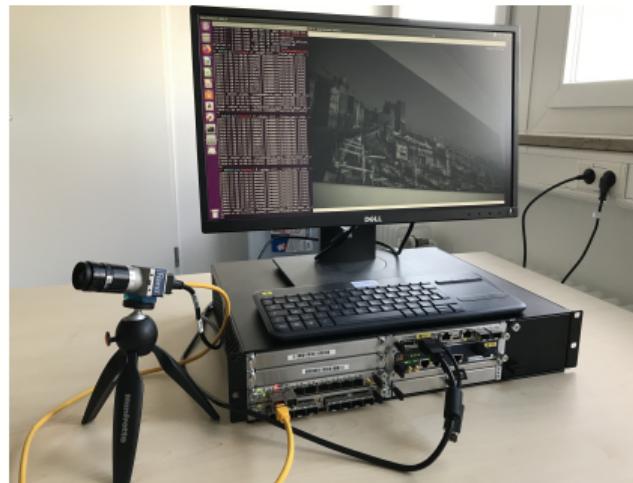
7th IBIC, Shanghai 2018

microTCA TECHNOLOGY LAB **HELMHOLTZ** RESEARCH FOR GRAND CHALLENGES

from http://ibic2018.vrws.de/talks/thoa02_talk.pdf



- ▶ DIPC-7050 is an IP stack for interfacing GigE Vision cameras with a FPGA
- ▶ Devices are running with up to (10)Gb Ethernet

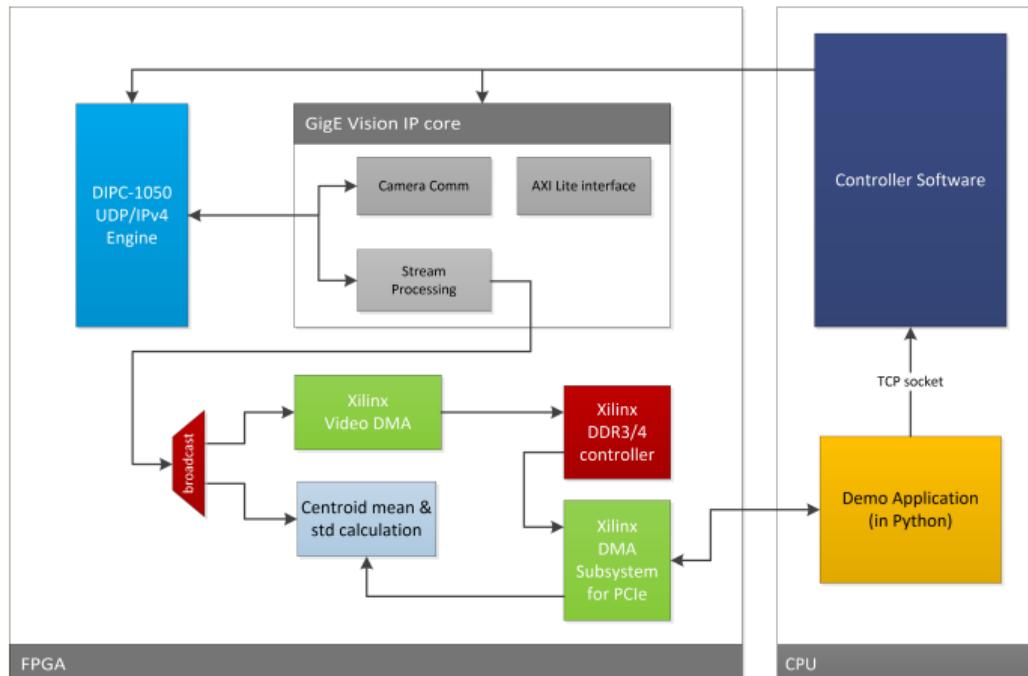


Talk on Thursday at 11:45:

"Certification and improvements of MicroTCA Technology Lab's GigE Vision Stack"

Example 2: GigE Vision camera interfacing

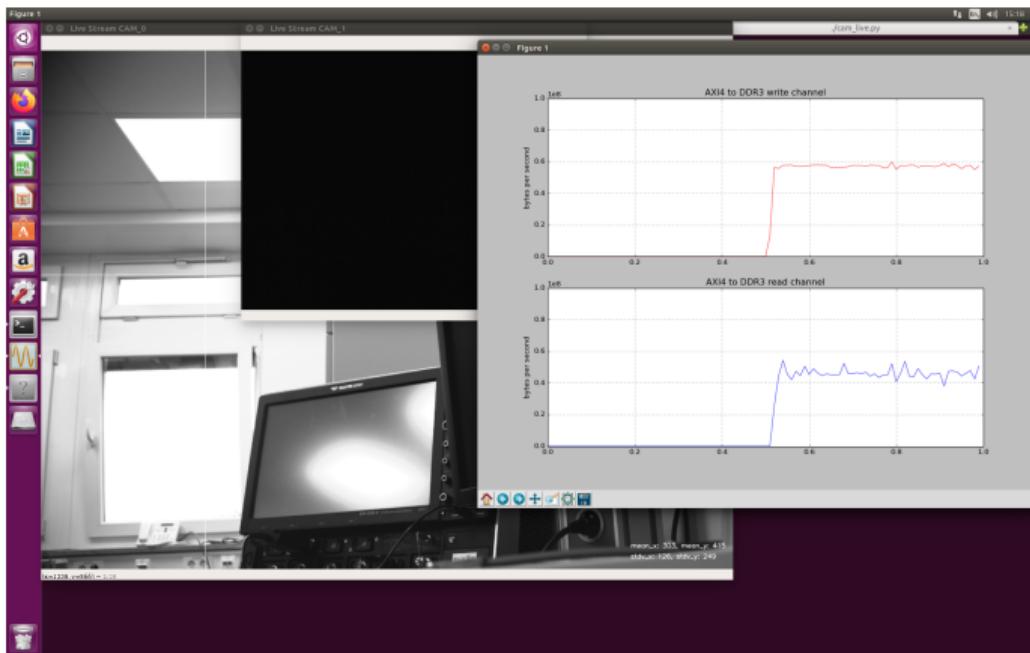
DMA over PCIe with FPGA
Stubbe, Marjanović, 2019-12-03 Page 15/70



DIPC-7050 GigE Vision Stack base configuration

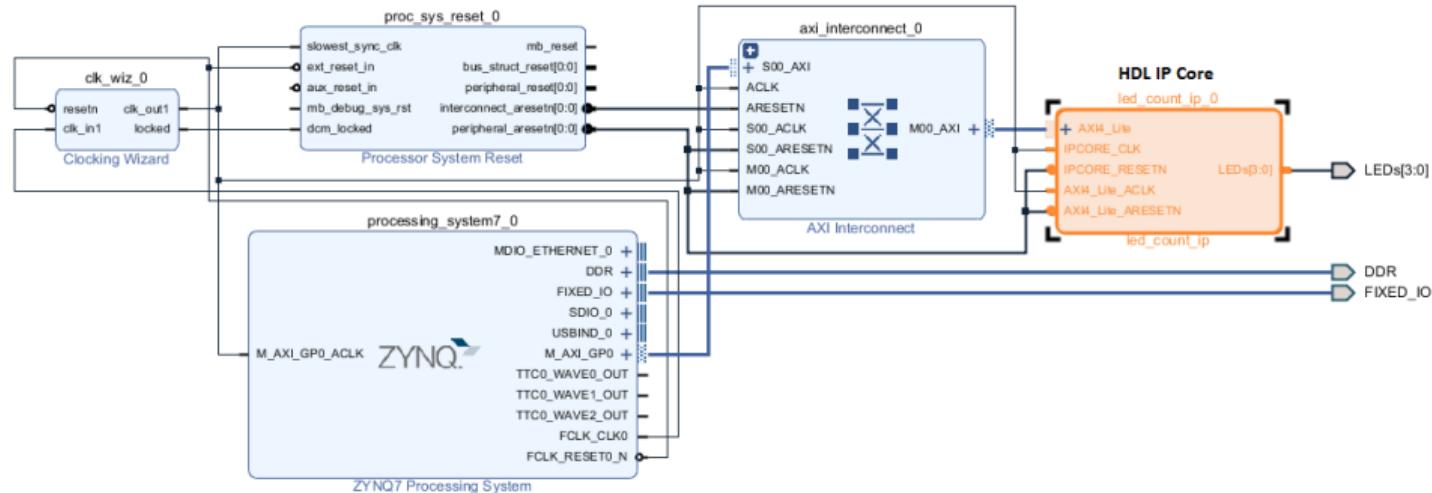
Example 2: GigE Vision camera interfacing

DMA over PCIe with FPGA
Stubbe, Marjanović, 2019-12-03 Page 16/70



Running DIPC-7050 demo application with AXI performance monitor

Basics



from <https://www.mathworks.com/help/hdlcoder/examples/define-and-register-custom-board-and-reference-design-for-zynq-workflow.html>

- ▶ (Almost) all Xilinx IPs use AMBA AXI protocols for interfacing between each other
- ▶ Provide good interface for testing

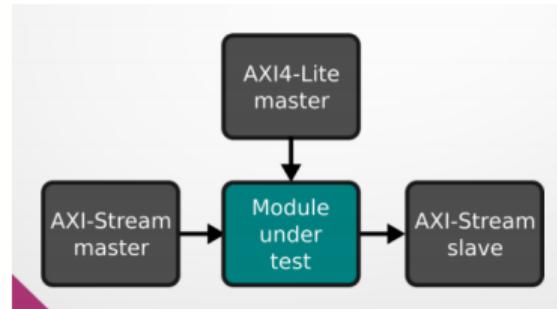


image from:

<https://indico.desy.de/indico/event/15974/session/6/contribution/15/material/slides/0.pdf>

another example (with Vivado):

https://github.com/j-marjanovic/chisel-stuff/tree/master/example-3-vivado-ip/vivado_tb_project

- ▶ Save documentation and communication efforts

Very common in Xilinx IPs for data streams

E.g.: FFT, 1/10/40/100 Gigabit Ethernet, high-perf PCIe

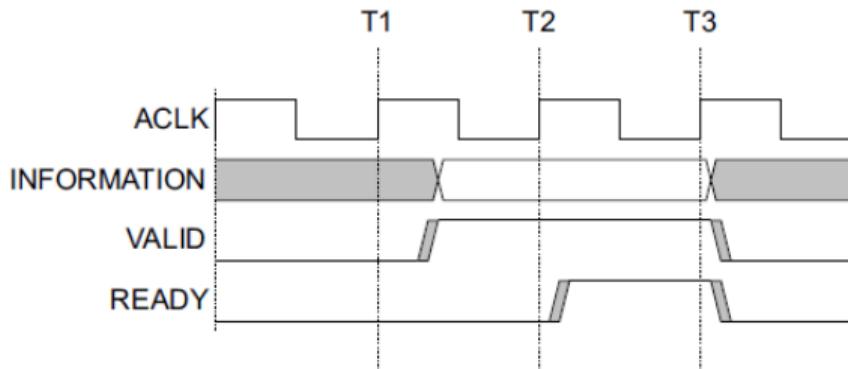


Figure A3-2 VALID before READY handshake

from AMBA® AXI™ and ACE™ Protocol Specification, available at

<https://www.arm.com/products/silicon-ip-system/embedded-system-design/amba-specifications>

Very common in Xilinx IPs for access to control/status registers

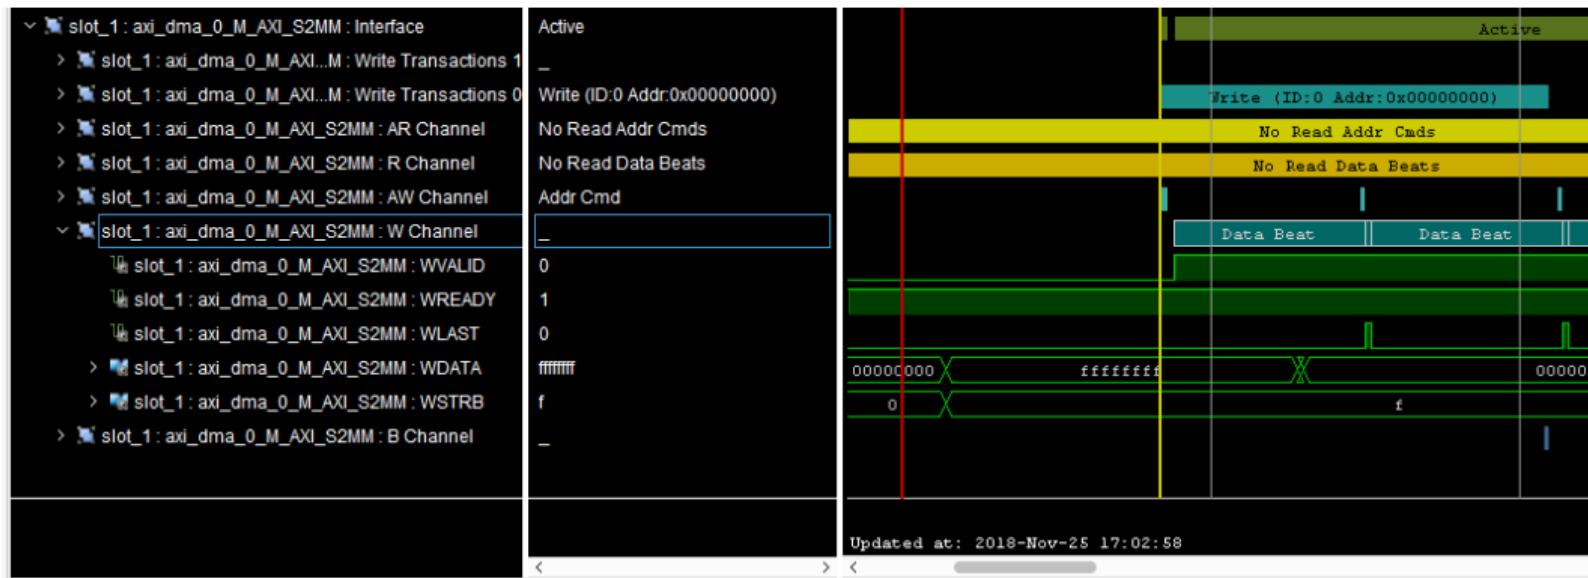
Composed of 5 independent channels:

- ▶ read address ($M \rightarrow S$)
- ▶ read data ($S \rightarrow M$)
- ▶ write address ($M \rightarrow S$)
- ▶ write data ($M \rightarrow S$)
- ▶ write response ($S \rightarrow M$)

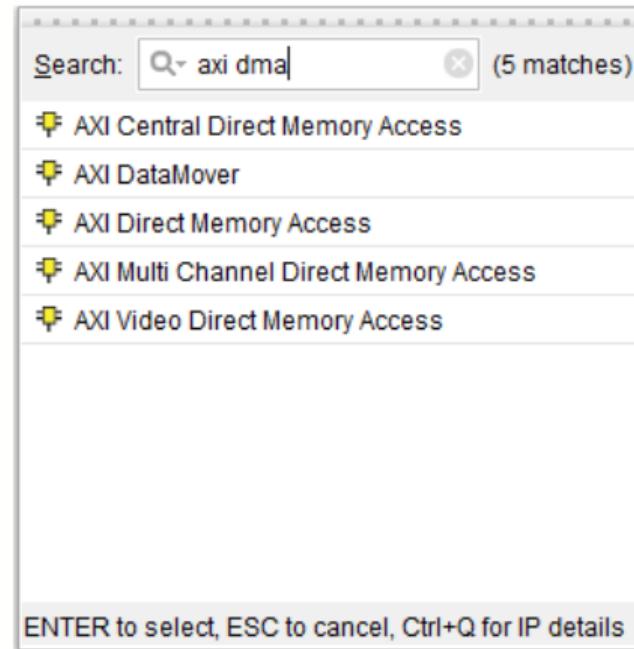
Three ways to implement AXI4-Lite interface:

- ▶ Vivado Wizard (Tools/Create and Package New IP)
- ▶ write the interface yourself (2 state machines)
- ▶ Vivado HLS (in C and C++, bundle all inputs/outputs into one interface)

Similar to AXI4-Lite, but support burst transfers (typically 16, 32 data beats per burst).
Very common in Xilinx IP for high-throughput interfaces (DDR3/DDR4, PCIe DMA, DMA)



Vivado provides a large variety of AXI DMA:



► Altera mSGDMA

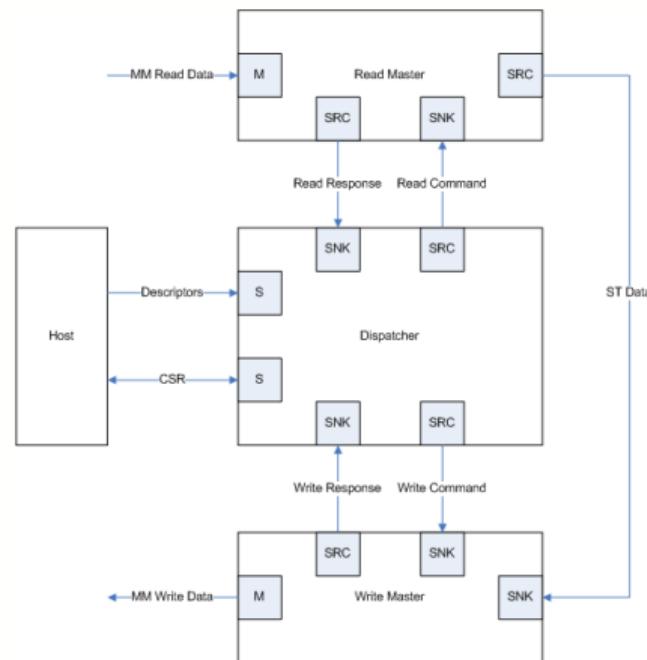
https://fpgawiki.intel.com/wiki/Modular_SGDMA

- IMHO the nicest SW interface
- uses FIFOs for commands and responses

► home-made DMA

- necessary in Virtex-5 times
- weeks to implements
- weeks to test all edge cases

► with Vivado HLS



Not really needed for this talk, but nice to understand what is going on behind the scenes.

Xilinx XDMA takes AXI transactions and creates PCIe TLP packets; these are passed to the Xilinx PCIe IP.

Figure 2-1: Non-Posted Read Transaction Protocol

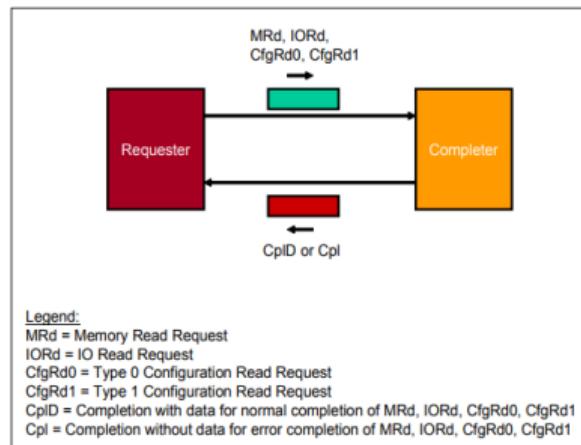
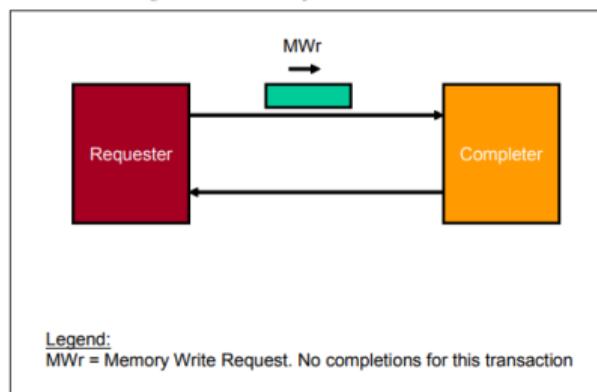


Figure 2-4: Posted Memory Write Transaction Protocol



from <https://www.mindshare.com/files/ebooks/PCI%20Express%20System%20Architecture.pdf>

Demo

- Decide for the hardware and physically setup the crate
- Compile the firmware image and program the FPGA
- Check PCIe link state and load linux driver
- Configure the registers in FPGA firmware via PCIe access
- Run the user application and check the results

DAMC-TCK7 is commercially available at N.A.T.:

<https://www.nateurope.com/products/NAT-AMC-TCK7.html>



DAMC-TCK7 is set up together with a Concurrent AM G64 CPU here in the front:



- Decide for the hardware and physically setup the crate
- Compile the firmware image and program the FPGA
- Check PCIe link state and load linux driver
- Configure the registers in FPGA firmware via PCIe access
- Run the user application and check the results

Available at <https://github.com/MicroTCA-Tech-Lab/damc-tck7-fpga-bsp>

The screenshot shows a GitHub repository page for 'damc-tck7-fpga-bsp'. The page includes the MicroTCA Technology Lab logo, location in Hamburg, Germany, and a link to their website. The navigation bar shows 1 repository, 3 people, 0 teams, 0 projects, and settings. The main content area displays the pinned repository 'damc-tck7-fpga-bsp' with a brief description: 'Board Support Package for DAMC-TCK7'. It shows the repository is written in Tcl and has 2 stars.

MicroTCA TECHNOLOGY LAB

Hamburg, Germany <https://techlab.desy.de>

Repositories 1 People 3 Teams 0 Projects 0 Settings

Pinned repositories

damc-tck7-fpga-bsp

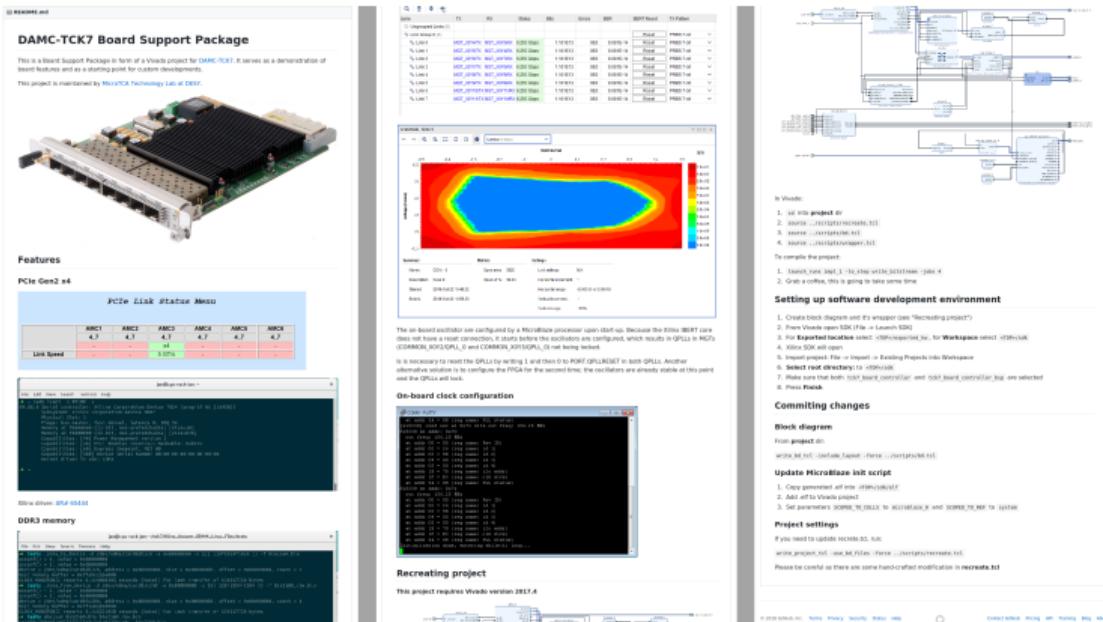
Board Support Package for DAMC-TCK7

Tcl ★ 2

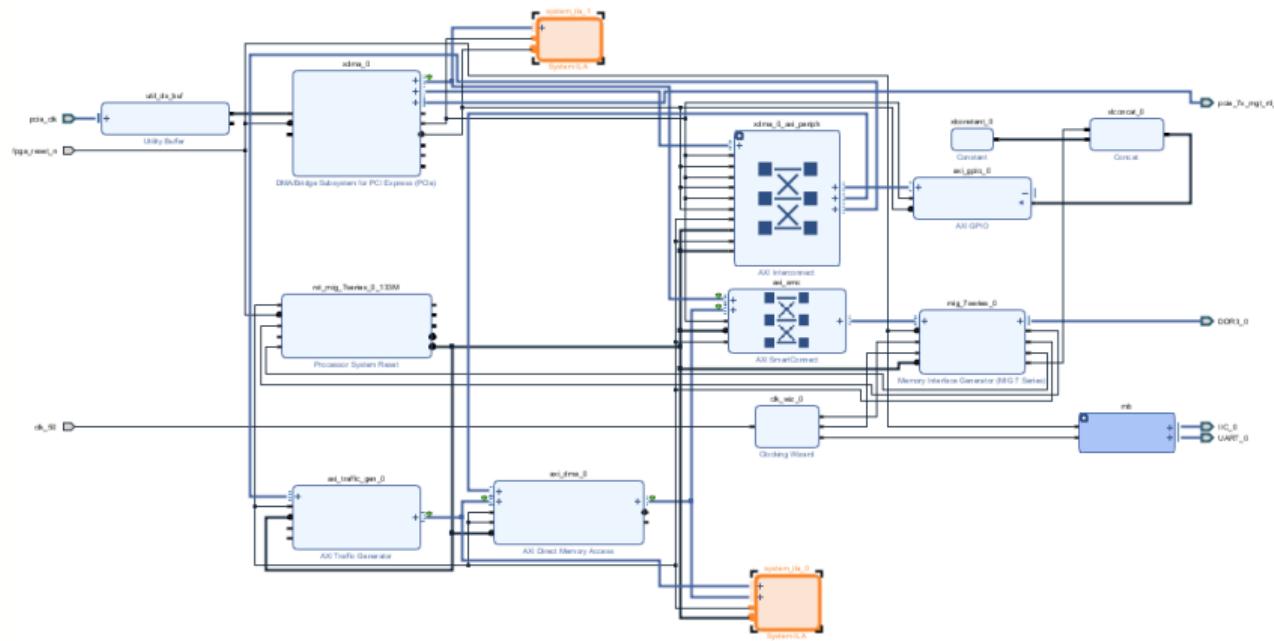
Start from the Board Support Package

DMA over PCIe with FPGA
Stubbe, Marjanović, 2019-12-03 Page 32/70

User can transfer the data from the CPU board to the DDR3 on DAMC-TCK7 board

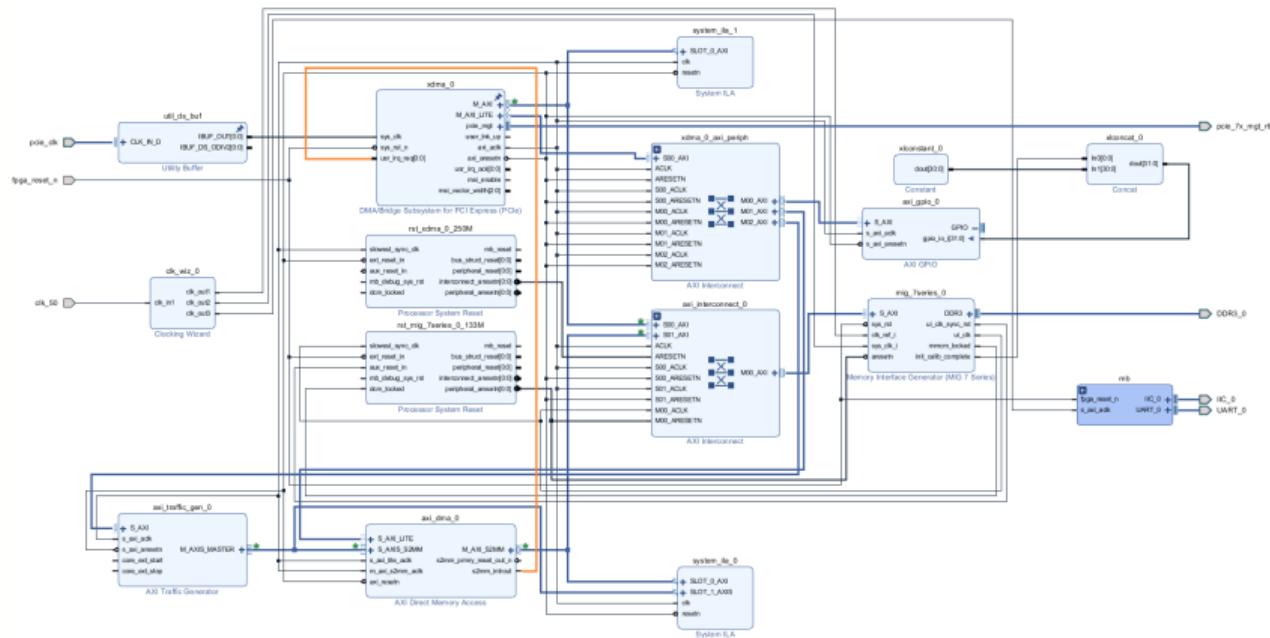


Vivado IP Integrator overview



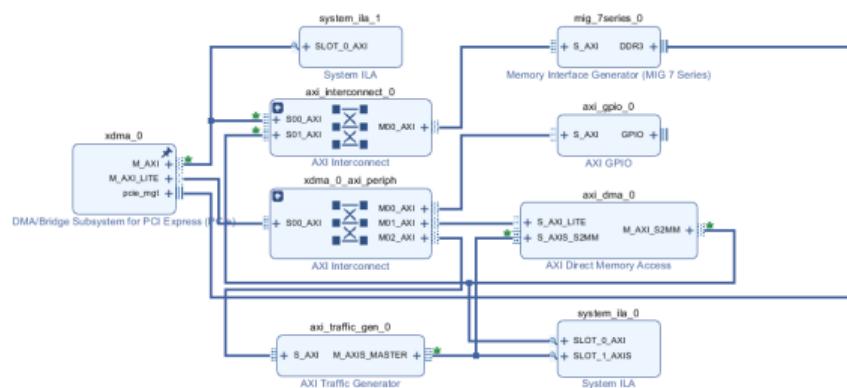
Two Integrated Logic Analyzers (ILAs) on important AXI connections are highlighted

Adding interrupt requests that will be monitored by user application



At this point our design includes:

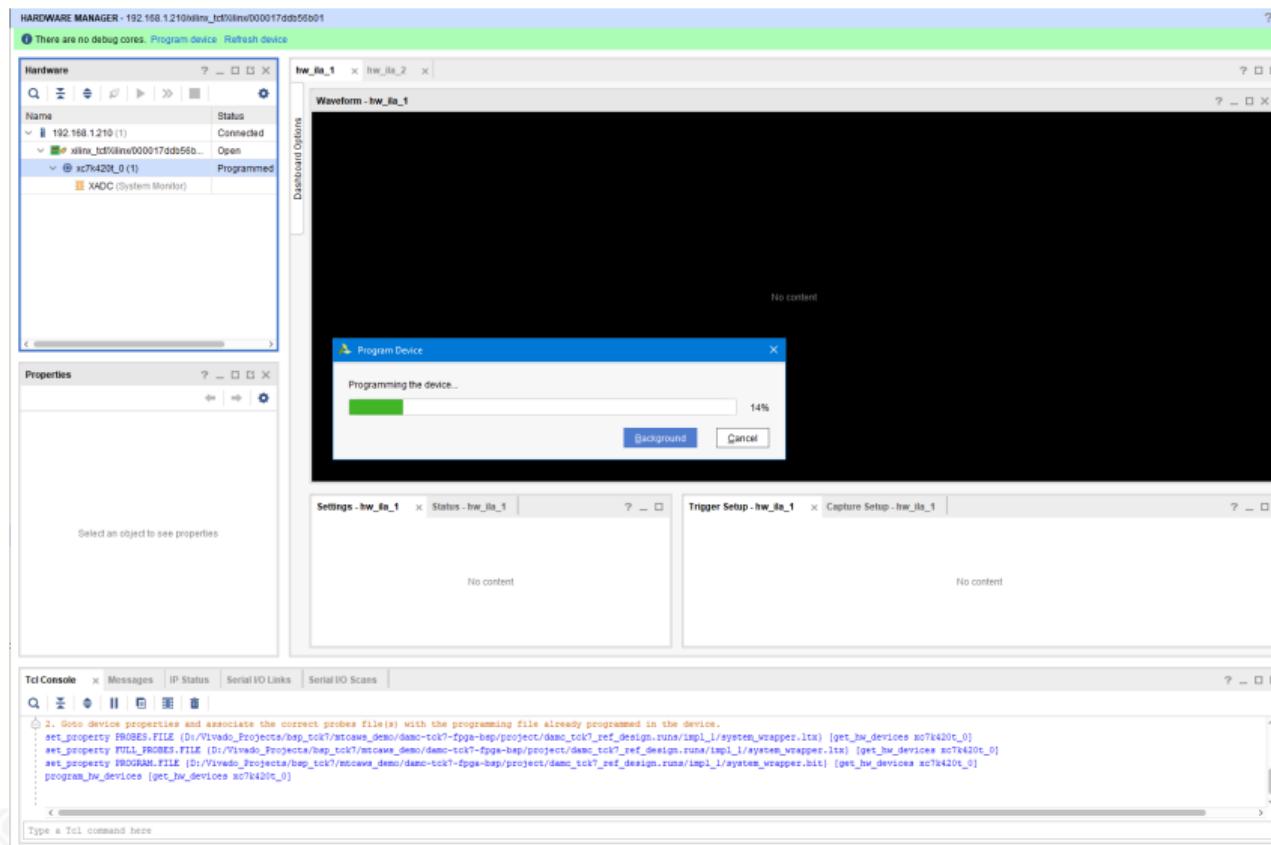
- ▶ DMA/Bridge Subsystem for PCI Express (from BSP)
- ▶ Memory Interface Generator (MIG) (from BSP)
- ▶ AXI DMA
- ▶ AXI Traffic Generator



Block design reduced to AXI connections

Program FPGA

DMA over PCIe with FPGA
Stubbe, Marjanović, 2019-12-03 Page 36/70



- Decide for the hardware and physically setup the crate
- Compile the firmware image and program the FPGA
- Check PCIe link state and load linux driver
- Configure the registers in FPGA firmware via PCIe access
- Run the user application and check the results

PCIe link states can be observed by accessing the MCH (telnet or webinterface):

NAT-MCH by N.A.T.

PCIe Link Status Menu								
	AMC1	AMC2	AMC3	AMC4	AMC5	AMC6	OPT1	RTM
Link Speed	-	-	x4	x4	x4	-	-	-

Linux driver provided by Xilinx (AR# 65444) needs to be downloaded and compiled.
Sources are available at https://github.com/Xilinx/dma_ip_drivers.

Xilinx QDMA IP Drivers https://xilinx.github.io/dma_ip_drivers/

28 commits 3 branches 0 packages 0 releases 4 contributors

Branch: master ▾ New pull request Create new file Upload files Find file Clone or download ▾

pankajdarak-xlnx Link QDMA IP drivers documentation to https://xilinx.github.io/dma_ip... ...	Latest commit 9505aa0 25 days ago
QDMA QDMA Linux reference driver 2019.1 Patch release	4 months ago
XDMA/linux-kernel need to check for non_incr_addr in both poll & isr mode	3 months ago
XVSEC/linux-kernel XVSEC driver: initial release	10 months ago
README.md Link QDMA IP drivers documentation to https://xilinx.github.io/dma_ip... ...	25 days ago

Get PCIe bus and device informations via `lspci`:

```
> sudo lspci -s 06:00 -v
06:00.0 Serial controller: Xilinx Corporation Device 7024 (prog-if 01 [16450])
    Subsystem: Xilinx Corporation Device 0007
    Physical Slot: 3
    Flags: bus master, fast devsel, latency 0, IRQ 212
    Memory at 9ac00000 (32-bit, non-prefetchable) [size=1M]
    Memory at 9ad00000 (32-bit, non-prefetchable) [size=64K]
    Capabilities: [40] Power Management version 3
    Capabilities: [48] MSI: Enable+ Count=1/1 Maskable- 64bit+
    Capabilities: [60] Express Endpoint, MSI 00
    Capabilities: [100] Device Serial Number 00-00-00-00-00-00-00-00
    Kernel driver in use: xdma
```

On the CPU side , Xilinx DMA driver provides several char devices:

```
$ ll /dev | grep xdma
drwxr-xr-x  3 root root          60 Jun 20 17:19 xdma
crw-rw-rw-  1 root root        238,  36 Jun 20 17:19 xdma0_c2h_0
crw-rw-rw-  1 root root        238,   1 Jun 20 17:19 xdma0_control
crw-rw-rw-  1 root root        238,  10 Jun 20 17:19 xdma0_events_0
crw-rw-rw-  1 root root        238,  11 Jun 20 17:19 xdma0_events_1
[...]
crw-rw-rw-  1 root root        238,  24 Jun 20 17:19 xdma0_events_14
crw-rw-rw-  1 root root        238,  25 Jun 20 17:19 xdma0_events_15
crw-rw-rw-  1 root root        238,  32 Jun 20 17:19 xdma0_h2c_0
crw-rw-rw-  1 root root        238,   0 Jun 20 17:19 xdma0_user
crw-rw-rw-  1 root root        238,   2 Jun 20 17:19 xdma0_xvc
```

On the CPU side , Xilinx DMA driver provides several **c2h** devices:

```
$ ll /dev | grep xdma
```

Mode	Major	Minor	Owner	Size	Date	Time	Type	Name
drwxr-xr-x	3	root root			60 Jun 20	17:19	xdma	
crw-rw-rw-	1	root root		238,	36 Jun 20	17:19	xdma0_c2h_0	
crw-rw-rw-	1	root root		238,	1 Jun 20	17:19	xdma0_control	
crw-rw-rw-	1	root root		238,	10 Jun 20	17:19	xdma0_events_0	
crw-rw-rw-	1	root root		238,	11 Jun 20	17:19	xdma0_events_1	
[...]								
crw-rw-rw-	1	root root		238,	24 Jun 20	17:19	xdma0_events_14	
crw-rw-rw-	1	root root		238,	25 Jun 20	17:19	xdma0_events_15	
crw-rw-rw-	1	root root		238,	32 Jun 20	17:19	xdma0_h2c_0	
crw-rw-rw-	1	root root		238,	0 Jun 20	17:19	xdma0_user	
crw-rw-rw-	1	root root		238,	2 Jun 20	17:19	xdma0_xvc	



- Decide for the hardware and physically setup the crate
- Compile the firmware image and program the FPGA
- Check PCIe link state and load linux driver
- Configure the registers in FPGA firmware via PCIe access
- Run the user application and check the results

To access the status and control registers (on AXI4-Lite interface), Python variant

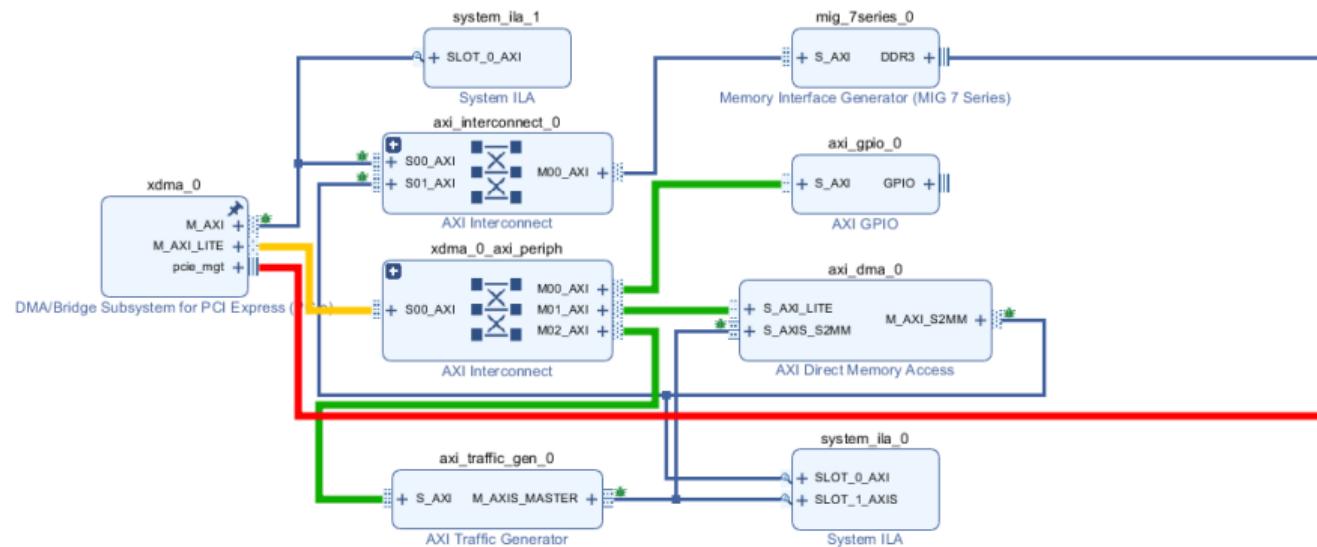
```
class HwAccessPcie(object):

    def __init__(self):
        user_filename = "/dev/xdma/card0/user"
        self.fd_user = os.open(user_filename, os.O_RDWR)
        self.mem = mmap.mmap(self.fd_user, 4 * 1024 * 1024)

    def __close__(self):
        self.mem.close()
        os.close(self.fd_user)

    def _rd32(self, addr):
        bs = self.mem[addr:addr + 4]
        return struct.unpack("I", bs)[0]

    def _wr32(self, addr, data):
        bs = struct.pack("I", data)
        self.mem[addr:addr + 4] = bs
```



Block design reduced to AXI connections

BAR0

- └─ 0x0000 AXI GPIO
 - └─ 0x0000 inputs
 - └─ bit[0] = DDR3 ready
- └─ 0x1000 AXI DMA
 - └─ 0x1030 S2MM_DMACR
 - └─ bit[0] = Run/Stop
 - └─ bit[1] = RSVD (const 1)
 - └─ 0x1034 S2MM_DMASR
 - └─ bit[0] = Halted
 - └─ bit[1] = Idle
 - └─ bit[11:4] = various errors
 - └─ 0x1048 S2MM_DA
 - └─ 0x104C S2MM_DA_MSB
 - └─ 0x1058 S2MM_LENGTH
- └─ 0x10000 AXI TPM

Use tool provided by Xilinx (`reg_rw`) to "peek" and "poke" registers.

1. Check DDR3 status

```
./reg_rw /dev/xdma/card0/user 0x0
```

Expected value: 0x1

2. Check if DMA is in idle

```
./reg_rw /dev/xdma/card0/user 0x1034
```

Expected value: 0x1

Get the data from AXI Traffic Generator into on-board DDR3 memory

1. Enable DMA

```
./reg_rw /dev/xdma/card0/user 0x1030 w 0x1
```

2. Check if DMA is not idle anymore

```
./reg_rw /dev/xdma/card0/user 0x1034
```

3. Set DMA destination address

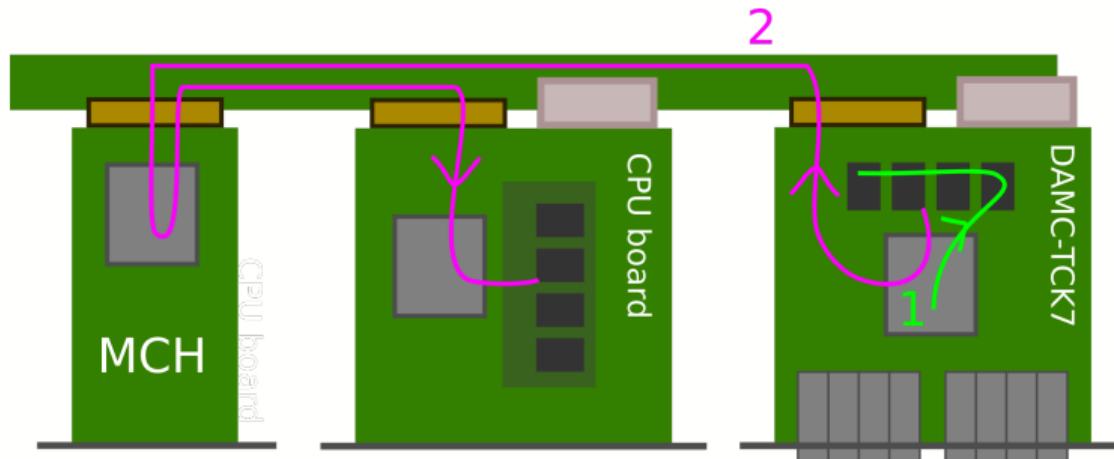
```
./reg_rw /dev/xdma/card0/user 0x1048 w 0x80000000
```

4. Start transfer

```
./reg_rw /dev/xdma/card0/user 0x1058 w 0x2000
```

- Decide for the hardware and physically setup the crate
- Compile the firmware image and program the FPGA
- Check PCIe link state and load linux driver
- Configure the registers in FPGA firmware via PCIe access
- Run the user application and check the results

1. AXI DMA will always be ready to accept data.
2. When an entire packet is received, AXI DMA raises interrupt request.
3. The user-space application monitors the interrupt request. When the AXI DMA has finished storing data into on-board DDR3, the user-space application is alerted.
4. The user-space application then copies the data from on-board DDR3 to its internal buffer.
5. The content of the buffer is compared to the SW implementation of LFSR.
6. A software script is used to re-trigger AXI Traffic Generator.



Content of App.cpp

```
/// implements LFSR as described in "AXI Traffic Generator v3.0"
class AxiTrafficGenLfsr;

/// interface to Xilinx DMA driver
class XdmaIsr;

/// general interface for register access (using wr32 and rd32 methods)
class XdmaRegAccess;

/// AXI DMA controller
class HwAxiDmaControl;

/// data consumer
class DataConsumer;

/// application capturing data and checking data integrity
int main();
```

We use the following script to trigger the AXI Traffic Generator.

```
#!/bin/bash
# script to configure and trigger AXI Traffic generator

AXI_TG_BASE=0x10000

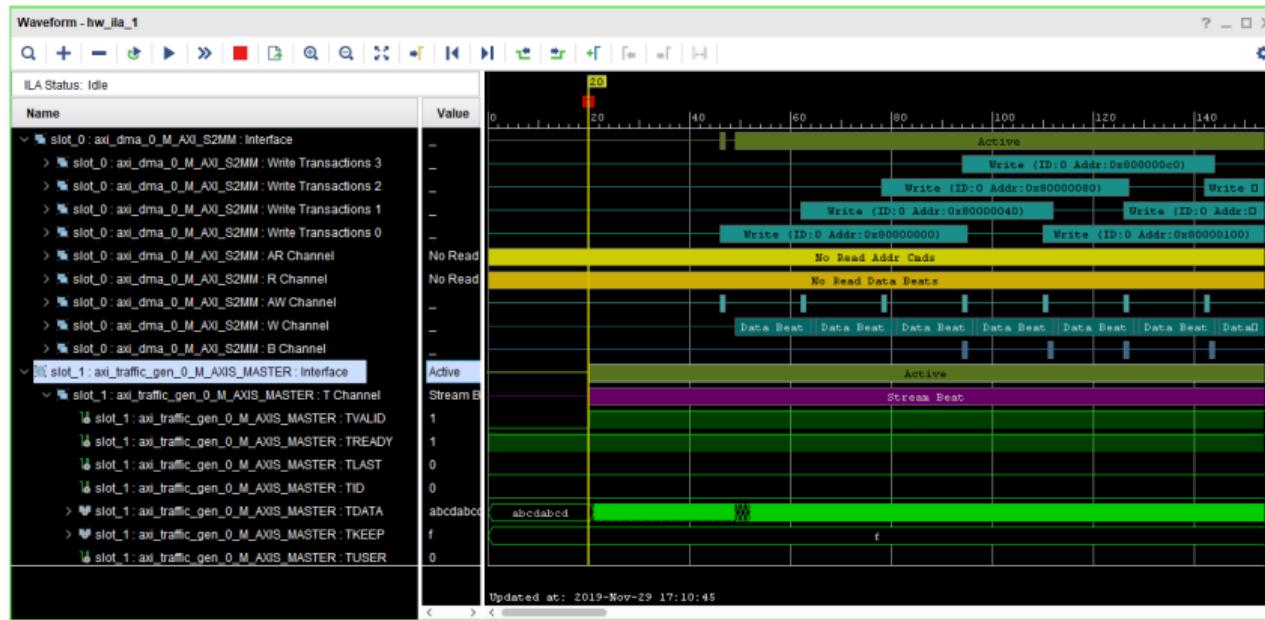
AXI_TG_ADDR_ST_CTRL=0x30
AXI_TG_ADDR_ST_CONF=0x34
AXI_TG_ADDR_TR_LEN=0x38

# set delay between packets
%reg_rw /dev/xdma/card0/user $((AXI_TG_BASE + AXI_TG_ADDR_ST_CONF)) w 0x10000

# set packet len, number of packets
reg_rw /dev/xdma/card0/user $((AXI_TG_BASE + AXI_TG_ADDR_TR_LEN)) w 0x101ff

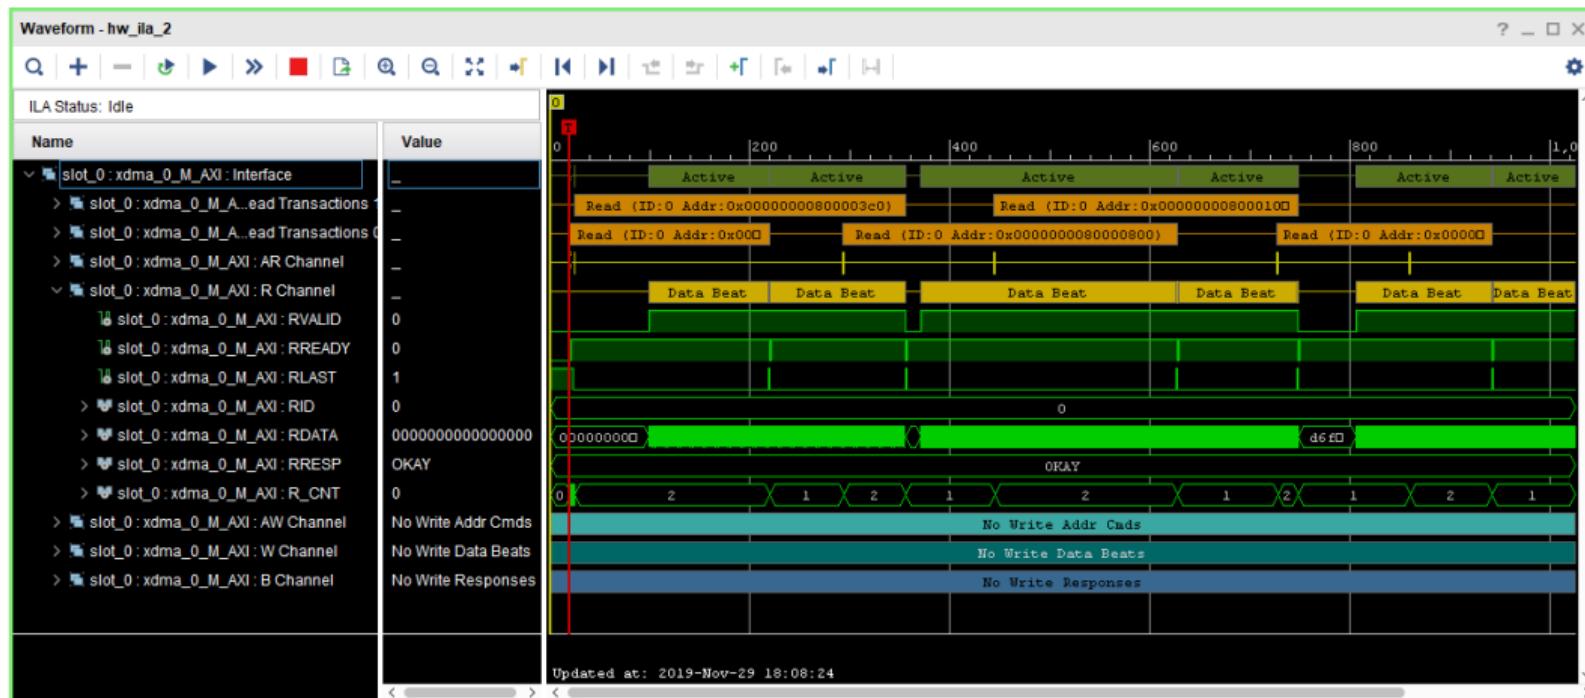
# trigger the traffic generation
reg_rw /dev/xdma/card0/user $((AXI_TG_BASE + AXI_TG_ADDR_ST_CTRL)) w 1
```

If we trigger ILA on AXI4-Stream VALID and READY being both '1' we are able to capture the DMA transfer into on-board DDR3 memory.



1. Data is now in on-board (DAMC-TCK7 board) DDR3 memory.
2. We need to use the second DMA to bring it into CPU.

Trigger on ARREADY and ARVALID to xdma/M_AXI:



(Note the incrementing address)

Output from application

```
./app
File Edit View Search Terminal Tabs Help
./app  x  jan@jan-Vi...  x  telnet MCH...  x  jan@jan-Vi...  x  jan@cpu-1...  x  +  ▾
→ app git:(master) X ./app
HwAxiDmaControl: status after run bit set = 0
XdmaIsr: received an event, value: 1
HwAxiDmaControl: status = 1002
HwAxiDmaControl: status = 2
DataConsumer::check, bytes_len = 2048
at  0 expect abcd, got: abcd | abcd |
at  1 expect d5e6, got: d5e6 | d5e6 |
at  2 expect eaf3, got: eaf3 | eaf3 |
at  3 expect f579, got: f579 | f579 |
at  4 expect 7abc, got: 7abc | 7abc |
at 1fc expect 33d0, got: 33d0 | 33d0 |
at 1fd expect 19e8, got: 19e8 | 19e8 |
at 1fe expect 8cf4, got: 8cf4 | 8cf4 |
at 1ff expect c67a, got: c67a | c67a |
DataConsumer: check, ok = 400, tot = 400 (100%)
XdmaIsr: received an event, value: 1
HwAxiDmaControl: status = 1002
HwAxiDmaControl: status = 2
DataConsumer::check, bytes_len = 2048
at  0 expect e33d, got: e33d | e33d |
at  1 expect f19e, got: f19e | f19e |
at  2 expect 78cf, got: 78cf | 78cf |
at  3 expect bc67, got: bc67 | bc67 |
at  4 expect 5e33, got: 5e33 | 5e33 |
at 1fc expect 2eb5, got: 2eb5 | 2eb5 |
at 1fd expect 175a, got: 175a | 175a |
at 1fe expect bad, got: bad | bad |
at 1ff expect 85d6, got: 85d6 | 85d6 |
DataConsumer: check, ok = 400, tot = 400 (100%)
```

Conclusion

- ▶ Xilinx provides a lot of infrastructure to help start the development
- ▶ MTCA Tech Lab is committed to provide open-source Board Support Packages for new boards (DAMC-TCK7, Zynq MPSOC FMC+ carrier, ...)
- ▶ Linux drivers are also provided by Xilinx
(<https://www.xilinx.com/support/answers/65444.html>)
Servers are a huge market and require better support for drivers/interfaces. Xilinx DMA provides everything out of the box.



TRANSFER MTCA TO RESEARCH AND INDUSTRY

- ▶ Custom developments
- ▶ High-end test & measurement services
- ▶ System configuration & integration
- ▶ LLRF design

MicroTCA training courses:

Next dates (see <https://techlab.desy.de/services/training/> for more):

- ▶ **Basic:** January 15-16, April 7-8, ...
- ▶ **Advanced:** February 26-27, June 3-4, ...

Thank you

<https://techlab.desy.de>

Deutsches Elektronen-Synchrotron DESY
A Research Centre of the Helmholtz Association
Notkestr. 85, 22607 Hamburg, Germany

Backup slides

Before we start with example, let's check if everything is OK with the system:

- FPGA programmed
- PCIe link is up (**show_link_state**, **show_ekey**, front-panel LED)
- PCIe device enumerated (**lspci**)
- Driver loaded (**lspci -v**)

```
system_i/xdma_0/inst/system_xdma_0_0_pcie2_to_pcie3_wrapper_i/
  pcie2_ip_i/s_axis_tx_tdata[*]
  pcie2_ip_i/s_axis_tx_tuser[*]
  pcie2_ip_i/s_axis_tx_tkeep[*]
  pcie2_ip_i/s_axis_tx_tvalid
  pcie2_ip_i/s_axis_tx_tready
  pcie2_ip_i/s_axis_tx_tlast
  pcie2_ip_i/m_axis_rx_tdata[*]
  pcie2_ip_i/m_axis_rx_tuser[*]
  pcie2_ip_i/m_axis_rx_tlast
  pcie2_ip_i/m_axis_rx_tready
  pcie2_ip_i/m_axis_rx_tvalid
```

Inspired by [drivers/net/wireless/broadcom/brcm80211/brcmfmac/pcie.c#L486](https://github.com/broadcom/linux-brcm80211/blob/master/brcmfmac/pcie.c#L486)

```
static void xdma_memcpy_from_dev (
    struct xdma_dev *lro,
    u32 offs,
    void *dstaddr,
    u32 len
) {
    void __iomem *address = lro->bar[lro->user_bar_idx] + offs;
    u32* dst32;

    len = len / 4;
    dst32 = (u32*)dstaddr;

    while (len) {
        *dst32 = ioread32(address);
        address += 4;
        dst32++;
        len--;
    }
}
```

see also <https://www.kernel.org/doc/html/v4.15/driver-api/device-io.html>

```
long char_ctrl_ioctl(struct file *filp, unsigned int cmd, unsigned long arg) {
    struct xdma_dev *lro;
    struct xdma_char *lro_char = (struct xdma_char *)filp->private_data;
    u64 t0_ns, t1_ns, duration_ns;

    BUG_ON(!lro_char);
    BUG_ON(lro_char->magic != MAGIC_CHAR);
    lro = lro_char->lro;
    BUG_ON(!lro);
    BUG_ON(lro->magic != MAGIC_DEVICE);

    printk(KERN_DEBUG "JAN: memcpy started, size: %lu\n", sizeof(tmp));
    t0_ns = ktime_get_ns();
    xdma_memcpy_from_dev(lro, USR_DDR3_OFFS, (void*)tmp, sizeof(tmp));
    t1_ns = ktime_get_ns();
    printk(KERN_DEBUG "JAN: memcpy done\n");
    duration_ns = t1_ns - t0_ns;
    printk(KERN_DEBUG "JAN: duration %llu ns (t1 = %llu, t0 = %llu)\n",
           duration_ns, t1_ns, t0_ns);
    return 0;
}
```

On Concurrent Tech AM900x
kernel 4.4.0-139-generic
Intel(R) Core(TM) i7-3612QE CPU @ 2.10GHz

```
[55775.132977] JAN: memcpy started, size: 1048576
[55775.498411] JAN: memcpy done
[55775.498414] JAN: duration 560049231 ns (t1 = 111312945419962, t0 = 111312385370731)

[55793.478833] JAN: memcpy started, size: 1048576
[55793.844333] JAN: memcpy done
[55793.844335] JAN: duration 468041144 ns (t1 = 111352256875650, t0 = 111351788834506)

[55819.083870] JAN: memcpy started, size: 1048576
[55819.449562] JAN: memcpy done
[55819.449564] JAN: duration 520045715 ns (t1 = 111401605213633, t0 = 111401085167918)
```

On ADLINK AMC-1000
kernel 4.15.0-39-generic
Intel(R) Core(TM)2 Duo CPU L7400 @ 1.50GHz

```
[ 357.333248] JAN: memcpy started, size: 1048576
[ 357.792152] JAN: memcpy done
[ 357.792157] JAN: duration 458898891 ns (t1 = 357792150659, t0 = 357333251768)

[ 360.061151] JAN: memcpy started, size: 1048576
[ 360.520068] JAN: memcpy done
[ 360.520073] JAN: duration 458910762 ns (t1 = 360520066218, t0 = 360061155456)

[ 362.647225] JAN: memcpy started, size: 1048576
[ 363.106262] JAN: memcpy done
[ 363.106268] JAN: duration 459030051 ns (t1 = 363106259232, t0 = 362647229181)
```

On external CPU

kernel 4.15.0-39-generic

Intel(R) Core(TM) i5-2500 CPU @ 3.30GHz

```
[ 225.678066] JAN: memcpy started, size: 1048576
[ 226.128774] JAN: memcpy done
[ 226.128775] JAN: duration 450702965 ns (t1 = 225831915216, t0 = 225381212251)

[ 229.493646] JAN: memcpy started, size: 1048576
[ 229.944742] JAN: memcpy done
[ 229.944743] JAN: duration 451090885 ns (t1 = 229647854461, t0 = 229196763576)

[ 232.285649] JAN: memcpy started, size: 1048576
[ 232.736626] JAN: memcpy done
[ 232.736627] JAN: duration 450971647 ns (t1 = 232439716846, t0 = 231988745199)
```

We inspect the data transferred with `dma_from_device` and note the pattern which was generated by AXI Traffic Generator (2x16-bit values).

```
jan@cpu-110008: ~/dma_demo/driver/Xilinx_Answer_65444_Linux_Files/tests
File Edit View Search Terminal Help
→ tests hexdump -C jan.bin
00000000 cd ab cd ab e6 d5 e6 d5 f3 ea f3 ea 79 f5 79 f5 | .....y.y.|  
00000010 bc 7a bc 7a 5e bd 5e bd af 5e af 5e 57 af 57 af |.z.z^.^..^W.W.|  
00000020 ab d7 ab d7 d5 eb ea 75 ea 75 f5 3a f5 3a |.....u.u.:.|  
00000030 7a 9d 7a 9d bd 4e bd 4e 5e a7 5e a7 af d3 af d3 |z.z..N.N^.^...|  
00000040 d7 e9 d7 e9 eb f4 eb f4 75 fa 75 fa 3a fd 3a fd |.....u.u.:.|  
00000050 9d 7e 9d 7e 4e 3f 4e 3f a7 1f a7 1f d3 0f d3 0f |.~.~N?N?....|  
00000060 e9 87 e9 87 f4 c3 f4 c3 fa e1 fa e1 fd f0 fd f0 |.....|  
00000070 7e 78 7e 78 3f 3c 3c 1f 9e 1f 9e 0f cf 0f cf |~x-x?<?<.....|  
00000080 87 67 87 67 c3 b3 c3 b3 e1 59 e1 59 f0 ac f0 ac |.g.g....Y.Y....|  
00000090 78 d6 78 d6 3c eb 3c eb 9e 75 9e 75 cf 3a cf 3a |x.x.<.<..u.u.:.|  
000000a0 67 9d 67 9d b3 4e b3 4e 59 a7 59 a7 ac d3 ac d3 |g.g..N.NY.Y....|  
000000b0 d6 e9 d6 e9 eb 74 eb 74 75 ba 75 ba 3a dd 3a dd |.....t.tu.u.:.|  
000000c0 9d 6e 9d 6e 4e b7 4e b7 a7 5b a7 5b d3 2d d3 2d |.n.nN.N..[.|-.-|  
000000d0 e9 96 e9 96 74 4b 74 4b ba a5 ba a5 dd 2d dd 2d |....tKtK.....|  
000000e0 6e 69 6e 69 b7 b4 b7 b4 5b 5a 5b 5a 2d ad 2d ad |nini...[Z[Z-...|  
000000f0 96 d6 96 d6 4b eb 4b eb a5 75 a5 75 d2 ba d2 ba |....K.K..u.u....|  
00000100 69 dd 69 dd b4 6e b4 6e 5a b7 5a b7 ad 5b ad 5b |i.i..n.nZ.Z..[.|  
00000110 d6 2d d6 2d eb 16 eb 16 75 8b 75 8b ba 45 ba 45 |.~.~.u.u..E.E|  
00000120 dd a2 dd a2 6e d1 6e d1 b7 68 b7 68 5b b4 5b b4 |....n.n..h.h[.|  
00000130 2d da 2d da 16 6d 16 6d 8b 36 8b 36 45 9b 45 9b |....m.m.6.6E.E.|  
00000140 a2 cd a2 cd d1 66 d1 66 68 33 68 33 b4 99 b4 99 |....f.fh3h3....|  
00000150 da 4c da 4c 6d a6 6d a6 36 d3 36 d3 9b e9 9b e9 |.L.Lm.m.6.6....|  
00000160 cd 74 cd 74 66 3a 66 3a 33 9d 33 9d 99 4e 99 4e |.t.tf:f:3.3..N.N|
```

LFSR Implementation used for Random Generation

The below linear feed-back shift register (LFSR) is used for random address or data generation:

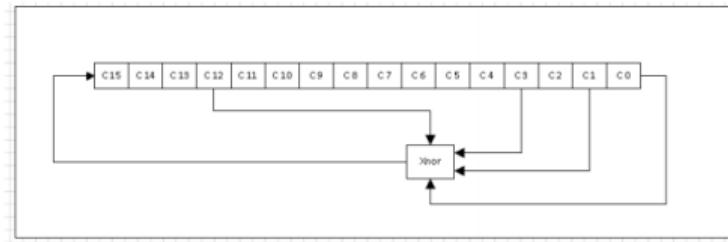


Figure 2-5: LFSR

Initially, the C0 to C15 flops are loaded with the input seed value. Later, it behaves as a shift register as per the architecture.

from Xilinx AXI Traffic Generator v3.0

https://www.xilinx.com/support/documentation/ip_documentation/axi_traffic_gen/v3_0/pg125-axi-traffic-gen.pdf